

customer segmentation

March 23, 2022

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import sea
from kneed import KneeLocator
from sklearn.datasets import make_blobs
from sklearn.cluster import KMeans
from sklearn.metrics import silhouette_score
from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA
from mpl_toolkits.mplot3d import Axes3D
```

Data is obtained from the following site: <https://www.kaggle.com/kaushiksuresh147/customer-segmentation>

0.0.1 Read Data and Make Sense of It

```
[2]: auto1 = pd.read_csv("/Users/owner/Desktop/DS_AI_projects/Customer segmentation_
↳project/Train.csv")
```

```
[3]: auto2 = pd.read_csv("/Users/owner/Desktop/DS_AI_projects/Customer segmentation_
↳project/Test.csv")
```

```
[4]: automobile = pd.concat([auto1, auto2])
automobile.head()
```

```
[4]:
```

	ID	Gender	Ever_Married	Age	Graduated	Profession	Work_Experience	\
0	462809	Male	No	22	No	Healthcare	1.0	
1	462643	Female	Yes	38	Yes	Engineer	NaN	
2	466315	Female	Yes	67	Yes	Engineer	1.0	
3	461735	Male	Yes	67	Yes	Lawyer	0.0	
4	462669	Female	Yes	40	Yes	Entertainment	NaN	

	Spending_Score	Family_Size	Var_1	Segmentation
0	Low	4.0	Cat_4	D
1	Average	3.0	Cat_4	A
2	Low	1.0	Cat_6	B

3	High	2.0	Cat_6	B
4	High	6.0	Cat_6	A

```
[5]: automobile.describe()
```

```
[5]:
```

	ID	Age	Work_Experience	Family_Size
count	10695.000000	10695.000000	9597.000000	10247.000000
mean	463468.088640	43.511828	2.619777	2.844052
std	2600.966411	16.774158	3.390790	1.536427
min	458982.000000	18.000000	0.000000	1.000000
25%	461220.500000	30.000000	0.000000	2.000000
50%	463451.000000	41.000000	1.000000	3.000000
75%	465733.500000	53.000000	4.000000	4.000000
max	467974.000000	89.000000	14.000000	9.000000

0.0.2 Imputation, Filling Missing Values, One Hot Encoding

```
[6]: automobile.columns = ['id', 'gender', 'married', 'age', 'graduated', '
    ↪ 'profession',
    'work_experience', 'spending_score', 'family_size', 'var_1',
    'segmentation']
```

```
[7]: automobile = automobile.drop(columns=["id", "segmentation"])
    automobile.head()
```

```
[7]:
```

	gender	married	age	graduated	profession	work_experience	\
0	Male	No	22	No	Healthcare	1.0	
1	Female	Yes	38	Yes	Engineer	NaN	
2	Female	Yes	67	Yes	Engineer	1.0	
3	Male	Yes	67	Yes	Lawyer	0.0	
4	Female	Yes	40	Yes	Entertainment	NaN	

	spending_score	family_size	var_1
0	Low	4.0	Cat_4
1	Average	3.0	Cat_4
2	Low	1.0	Cat_6
3	High	2.0	Cat_6
4	High	6.0	Cat_6

```
[8]: # treating numerical features of null values with mean values
    automobile.fillna(automobile[["age", "work_experience", "family_size"]].mean(),
    ↪ inplace=True)
```

```
[9]: for c in automobile.columns:
    a = automobile[c].value_counts()
    print(a)
```

```

Male      5841
Female    4854
Name: gender, dtype: int64
Yes       6163
No        4342
Name: married, dtype: int64
35        321
42        320
37        304
43        301
40        300
...
75         37
78         36
76         36
80         33
85         32
Name: age, Length: 67, dtype: int64
Yes       6570
No        4023
Name: graduated, dtype: int64
Artist          3318
Healthcare      1750
Entertainment   1250
Engineer        935
Doctor          930
Lawyer          844
Executive       775
Marketing       403
Homemaker       328
Name: profession, dtype: int64
1.000000      3127
0.000000      3087
2.619777      1098
9.000000       613
8.000000       612
2.000000       373
4.000000       346
3.000000       337
5.000000       270
6.000000       265
7.000000       256
14.000000       66
10.000000       64
11.000000       64
12.000000       60
13.000000       57
Name: work_experience, dtype: int64

```

```

Low          6494
Average      2599
High         1602
Name: spending_score, dtype: int64
2.000000     3158
1.000000     1965
3.000000     1952
4.000000     1823
5.000000      812
2.844052      448
6.000000      290
7.000000      122
8.000000       65
9.000000       60
Name: family_size, dtype: int64
Cat_6        6910
Cat_4        1475
Cat_3        1089
Cat_2         563
Cat_7         269
Cat_1         167
Cat_5         114
Name: var_1, dtype: int64

```

```
[10]: automobile.isnull().sum()
```

```

[10]: gender          0
      married        190
      age            0
      graduated       102
      profession      162
      work_experience   0
      spending_score   0
      family_size      0
      var_1           108
      dtype: int64

```

```

[11]: # binary encoding for categorical features with only 2 possible values.

      # dict = {# some kind of dictionary with categorical keys mapping to numerical
      ↪ values}
      # df['feature'] = df['feature'].map(dict)

```

```

[12]: # one-hot encoding for categorical features with 2 or more possible values.
      automobiles = pd.get_dummies(automobile, prefix = ['profession'], columns =
      ↪ ['profession'])

```

```

automobiles = pd.get_dummies(automobiles, prefix = ['gender'], columns =
↳ ['gender'])
automobiles = pd.get_dummies(automobiles, prefix = ['married'], columns =
↳ ['married'])
automobiles = pd.get_dummies(automobiles, prefix = ['graduated'], columns =
↳ ['graduated'])
automobiles = pd.get_dummies(automobiles, prefix = ['var_1'], columns =
↳ ['var_1'])
automobiles.head(10)

```

```

[12]:  age  work_experience  spending_score  family_size  profession_Artist  \
0    22           1.000000           Low           4.0              0
1    38           2.619777       Average           3.0              0
2    67           1.000000           Low           1.0              0
3    67           0.000000           High           2.0              0
4    40           2.619777           High           6.0              0
5    56           0.000000       Average           2.0              1
6    32           1.000000           Low           3.0              0
7    33           1.000000           Low           3.0              0
8    61           0.000000           Low           3.0              0
9    55           1.000000       Average           4.0              1

```

```

      profession_Doctor  profession_Engineer  profession_Entertainment  \
0                    0                    0                    0
1                    0                    1                    0
2                    0                    1                    0
3                    0                    0                    0
4                    0                    0                    1
5                    0                    0                    0
6                    0                    0                    0
7                    0                    0                    0
8                    0                    1                    0
9                    0                    0                    0

```

```

      profession_Executive  profession_Healthcare  ...  married_Yes  \
0                    0                    1  ...              0
1                    0                    0  ...              1
2                    0                    0  ...              1
3                    0                    0  ...              1
4                    0                    0  ...              1
5                    0                    0  ...              1
6                    0                    1  ...              0
7                    0                    1  ...              0
8                    0                    0  ...              1
9                    0                    0  ...              1

```

```

      graduated_No  graduated_Yes  var_1_Cat_1  var_1_Cat_2  var_1_Cat_3  \

```

0	1	0	0	0	0
1	0	1	0	0	0
2	0	1	0	0	0
3	0	1	0	0	0
4	0	1	0	0	0
5	1	0	0	0	0
6	0	1	0	0	0
7	0	1	0	0	0
8	0	1	0	0	0
9	0	1	0	0	0

	var_1_Cat_4	var_1_Cat_5	var_1_Cat_6	var_1_Cat_7
0	1	0	0	0
1	1	0	0	0
2	0	0	1	0
3	0	0	1	0
4	0	0	1	0
5	0	0	1	0
6	0	0	1	0
7	0	0	1	0
8	0	0	0	1
9	0	0	1	0

[10 rows x 26 columns]

```
[13]: automobiles.drop(columns=["gender_Male"], inplace=True)
```

```
[14]: # Ordinal feature encoding for features that have ordered values
from sklearn.preprocessing import LabelEncoder

#label encoder can't handle missing values
automobiles['spending_score'] = automobiles['spending_score'].fillna('None')
# Label encode ord_1 feature
label_encoder = LabelEncoder()
automobiles['spending_score'] = label_encoder.
    ↳fit_transform(automobiles['spending_score'])
# Print sample of dataset
automobiles.head(10)
```

```
[14]:  age  work_experience  spending_score  family_size  profession_Artist  \
0    22         1.000000             2           4.0              0
1    38         2.619777             0           3.0              0
2    67         1.000000             2           1.0              0
3    67         0.000000             1           2.0              0
4    40         2.619777             1           6.0              0
5    56         0.000000             0           2.0              1
6    32         1.000000             2           3.0              0
```

7	33	1.000000	2	3.0	0
8	61	0.000000	2	3.0	0
9	55	1.000000	0	4.0	1

	profession_Doctor	profession_Engineer	profession_Entertainment	\
0	0	0	0	
1	0	1	0	
2	0	1	0	
3	0	0	0	
4	0	0	1	
5	0	0	0	
6	0	0	0	
7	0	0	0	
8	0	1	0	
9	0	0	0	

	profession_Executive	profession_Healthcare	...	married_Yes	\
0	0	1	...	0	
1	0	0	...	1	
2	0	0	...	1	
3	0	0	...	1	
4	0	0	...	1	
5	0	0	...	1	
6	0	1	...	0	
7	0	1	...	0	
8	0	0	...	1	
9	0	0	...	1	

	graduated_No	graduated_Yes	var_1_Cat_1	var_1_Cat_2	var_1_Cat_3	\
0	1	0	0	0	0	
1	0	1	0	0	0	
2	0	1	0	0	0	
3	0	1	0	0	0	
4	0	1	0	0	0	
5	1	0	0	0	0	
6	0	1	0	0	0	
7	0	1	0	0	0	
8	0	1	0	0	0	
9	0	1	0	0	0	

	var_1_Cat_4	var_1_Cat_5	var_1_Cat_6	var_1_Cat_7
0	1	0	0	0
1	1	0	0	0
2	0	0	1	0
3	0	0	1	0
4	0	0	1	0
5	0	0	1	0

6	0	0	1	0
7	0	0	1	0
8	0	0	0	1
9	0	0	1	0

[10 rows x 25 columns]

```
[15]: # feature standardization
scaler = StandardScaler().fit(automobiles)
scaled_features = scaler.transform(automobiles)
scaled_features = pd.DataFrame(scaled_features, columns = automobiles.columns)
scaled_features.head()
```

```
[15]:      age  work_experience  spending_score  family_size  profession_Artist \
0 -1.282499      -0.504312      0.750575      0.768669      -0.670654
1 -0.328606      0.000000     -1.610424      0.103701      -0.670654
2  1.400325     -0.504312      0.750575     -1.226237      -0.670654
3  1.400325     -0.815659     -0.429925     -0.561268      -0.670654
4 -0.209369      0.000000     -0.429925      2.098607      -0.670654

      profession_Doctor  profession_Engineer  profession_Entertainment \
0      -0.308607      -0.309514      -0.363793
1      -0.308607      3.230867      -0.363793
2      -0.308607      3.230867      -0.363793
3      -0.308607     -0.309514      -0.363793
4      -0.308607     -0.309514      2.748818

      profession_Executive  profession_Healthcare  ...  married_Yes \
0      -0.279508      2.260847  ...  -1.166141
1      -0.279508     -0.442312  ...   0.857529
2      -0.279508     -0.442312  ...   0.857529
3      -0.279508     -0.442312  ...   0.857529
4      -0.279508     -0.442312  ...   0.857529

      graduated_No  graduated_Yes  var_1_Cat_1  var_1_Cat_2  var_1_Cat_3 \
0      1.287814     -1.262033     -0.125946     -0.235726     -0.3367
1     -0.776510      0.792372     -0.125946     -0.235726     -0.3367
2     -0.776510      0.792372     -0.125946     -0.235726     -0.3367
3     -0.776510      0.792372     -0.125946     -0.235726     -0.3367
4     -0.776510      0.792372     -0.125946     -0.235726     -0.3367

      var_1_Cat_4  var_1_Cat_5  var_1_Cat_6  var_1_Cat_7
0      2.500169     -0.103798     -1.351158     -0.160627
1      2.500169     -0.103798     -1.351158     -0.160627
2     -0.399973     -0.103798      0.740106     -0.160627
3     -0.399973     -0.103798      0.740106     -0.160627
4     -0.399973     -0.103798      0.740106     -0.160627
```


[5 rows x 25 columns]


0.0.3 Make a Plot on the “Elbow Method”

```
[16]: # store the squared sum of distances from the mean of each cluster to the data_
      ↪points to that mean
sse = []

for cluster in range(1,10):
    kmeans = KMeans(n_clusters = cluster, init='k-means++')
    kmeans.fit(automobiles)
    sse.append(kmeans.inertia_)

# converting the results into a dataframe and plotting them
clustering = pd.DataFrame({'Cluster':range(1,10), 'SSE':sse})
#plt.figure(figsize=(12,6))
plt.plot(clustering['Cluster'], clustering['SSE'], marker='o')
plt.xlabel('Number of clusters')
plt.ylabel('Inertia')
```

```
[16]: Text(0, 0.5, 'Inertia')
```



```
[17]: # based on the elbow method, the best number of clusters to set would be 4,
# since after K=4, the inertia decreases in a linear manner.

# Build a model with 4 clusters
kmeans = KMeans(n_clusters=4, init='k-means++')
kmeans.fit(automobiles)

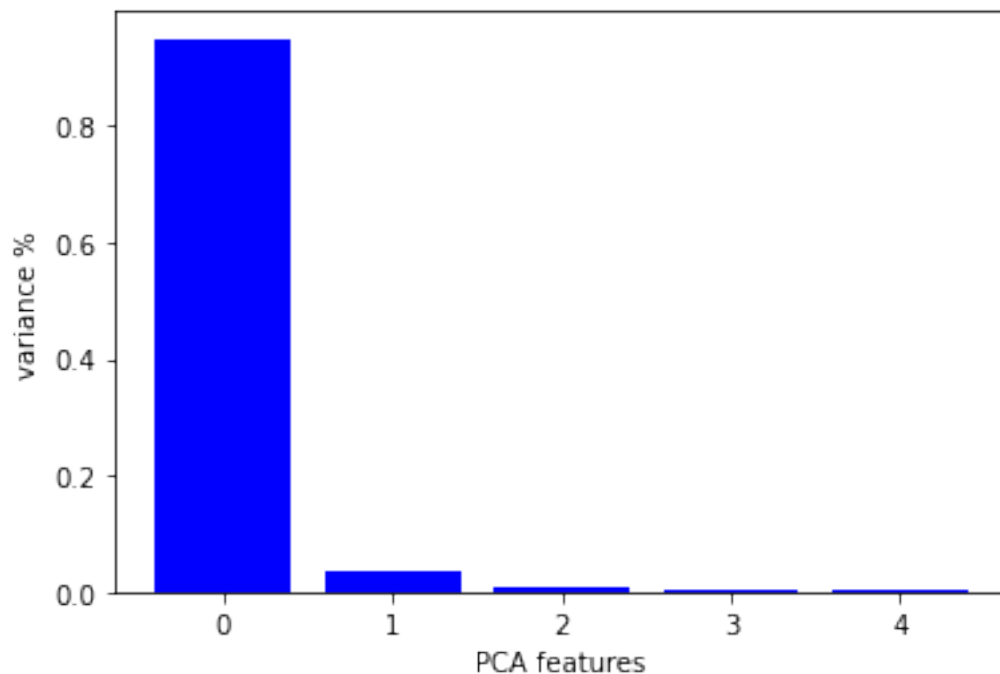
# Print the silhouette score for the model above
print(silhouette_score(automobiles, kmeans.labels_, metric='euclidean'))
```

0.42861278447347173

0.0.4 Try Improve the Quality of Clusters by Selecting Useful Features

```
[18]: # attempt to improve the model using PCA to reduce dimensionality.
pca = PCA(n_components=5)
pca_fit = pca.fit_transform(automobiles)
n_comp = range(pca.n_components_)
plt.bar(n_comp, pca.explained_variance_ratio_, color='blue')
plt.xlabel('PCA features')
plt.ylabel('variance %')

pca_components = pd.DataFrame(pca_fit)
```



```
[19]: pca_components.head()
```

```
[19]:
```

	0	1	2	3	4
0	-21.502903	-2.450373	0.521449	0.782306	-0.633519
1	-5.476260	-0.192459	0.297044	-1.566620	-0.290469
2	23.568221	-0.689226	-1.477698	0.566654	0.225764
3	23.602871	-1.765042	-0.394367	-0.085947	0.124108
4	-3.558898	-0.310606	2.998802	-0.374295	0.823571


```
[20]: # building a new model after PCA.

inertia = []

for cluster in range(1,10):
    kmeans = KMeans(n_clusters = cluster)
    kmeans.fit(pca_components.iloc[:, :2])
    inertia.append(kmeans.inertia_)

# converting the results into a dataframe and plotting them
clusters = pd.DataFrame({'Cluster':range(1,10), 'Inertia':inertia})
# plt.figure(figsize=(12,6))
plt.plot(clusters['Cluster'], clusters['Inertia'], marker='o')
plt.xlabel('Number of clusters')
plt.ylabel('Inertia')
```

```
[20]: Text(0, 0.5, 'Inertia')
```



```
[21]: # still, K=4 seems like the best approach.
kmeans_new = KMeans(n_clusters=4)
kmeans_new.fit(pca_components.iloc[:, :2])

# silhouette score
print(silhouette_score(pca_components.iloc[:, :2], kmeans_new.labels_,
    ↪metric='euclidean'))
```

0.47313464575183084


We see roughly 4.5% increase in the Silhouette score!

```
[22]: model = KMeans(n_clusters=4)

clusters = model.fit_predict(pca_components.iloc[:, :2])
automobiles["label"] = clusters

fig = plt.figure(figsize=(21,10))
ax = fig.add_subplot(111, projection='3d')
ax.scatter(automobiles.age[automobiles.label == 0],
    ↪automobiles["work_experience"][automobiles.label == 0],
    ↪automobiles["spending_score"][automobiles.label == 0], c='blue', s=60)
ax.scatter(automobiles.age[automobiles.label == 1],
    ↪automobiles["work_experience"][automobiles.label == 1],
    ↪automobiles["spending_score"][automobiles.label == 1], c='red', s=60)
ax.scatter(automobiles.age[automobiles.label == 2],
    ↪automobiles["work_experience"][automobiles.label == 2],
    ↪automobiles["spending_score"][automobiles.label == 2], c='green', s=60)
ax.scatter(automobiles.age[automobiles.label == 3],
    ↪automobiles["work_experience"][automobiles.label == 3],
    ↪automobiles["spending_score"][automobiles.label == 3], c='orange', s=60)

ax.view_init(30, 40)
plt.show()
```



```
[23]: automobiles.columns
```

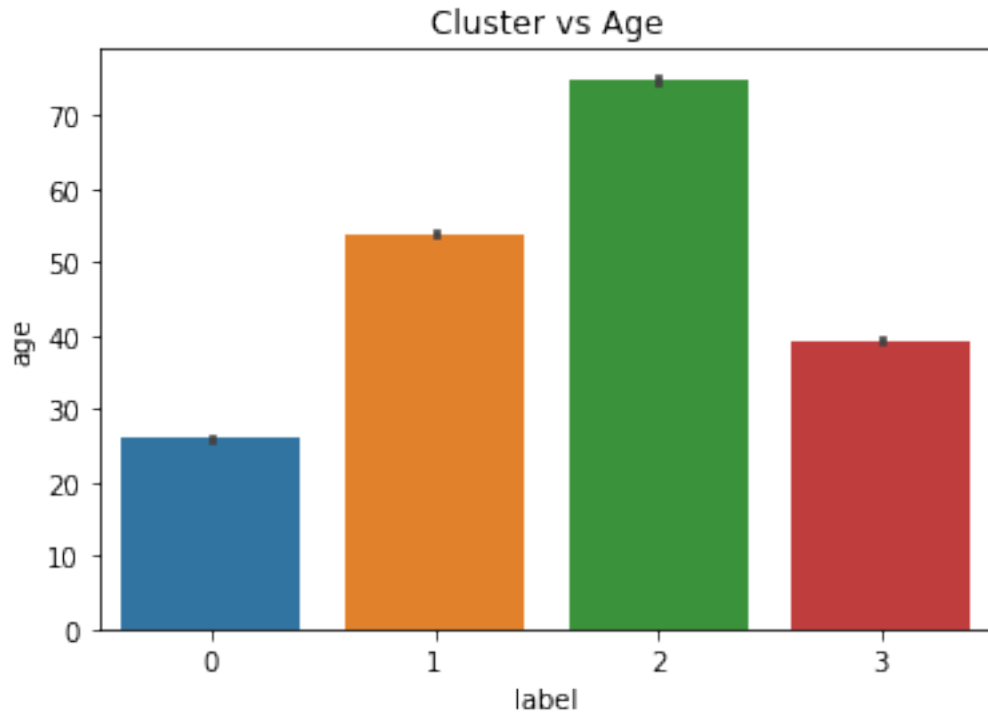
```
[23]: Index(['age', 'work_experience', 'spending_score', 'family_size',
          'profession_Artist', 'profession_Doctor', 'profession_Engineer',
          'profession_Entertainment', 'profession_Executive',
          'profession_Healthcare', 'profession_Homemaker', 'profession_Lawyer',
          'profession_Marketing', 'gender_Female', 'married_No', 'married_Yes',
          'graduated_No', 'graduated_Yes', 'var_1_Cat_1', 'var_1_Cat_2',
          'var_1_Cat_3', 'var_1_Cat_4', 'var_1_Cat_5', 'var_1_Cat_6',
          'var_1_Cat_7', 'label'],
          dtype='object')
```

As you can see, the clusters look like they can be separated quite easily, except a few spots where different colors overlap a bit. As far as the results go, I'm satisfied with the clusters!

0.0.5 Take a Look at the Demographics

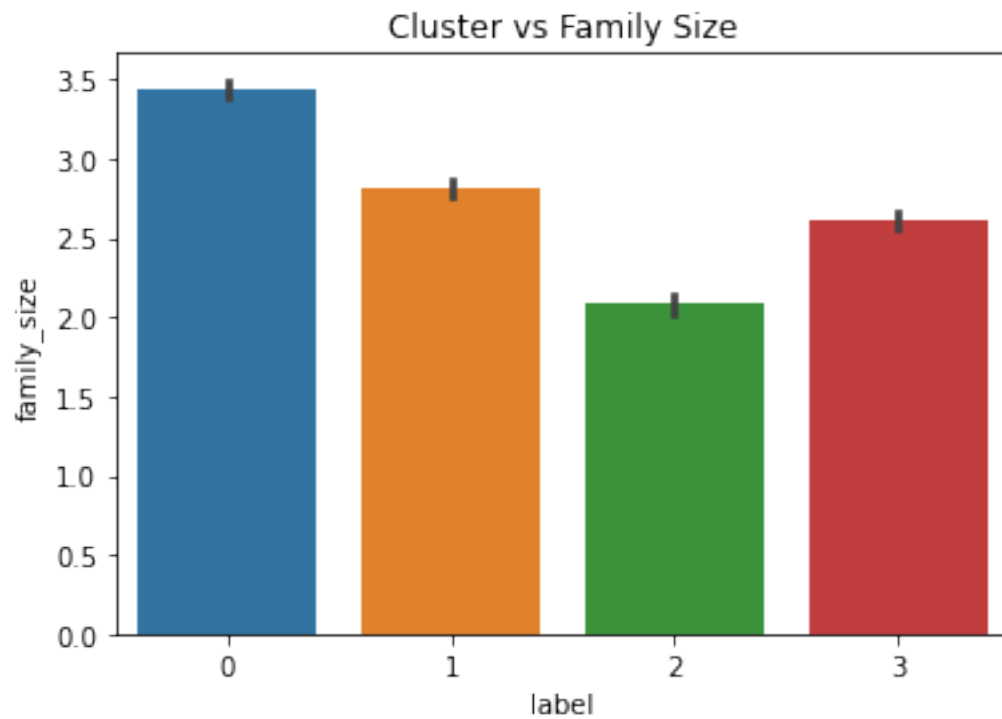
```
[24]: sns.barplot(x='label', y='age', data=automobiles).set(title='Cluster vs Age')
```

```
[24]: [Text(0.5, 1.0, 'Cluster vs Age')]
```



```
[25]: sns.barplot(x='label', y='family_size', data=automobiles).set(title='Cluster vs Family Size')
```

```
[25]: [Text(0.5, 1.0, 'Cluster vs Family Size')]
```



```
[26]: sns.barplot(x='label', y='spending_score', data=automobiles).set(title='Cluster_
      ↳vs Spending Score')
```

```
[26]: [Text(0.5, 1.0, 'Cluster vs Spending Score')]
```



```
[27]: sns.barplot(x='label', y='work_experience', data=automobiles).  
      ↪ set(title='Cluster vs Work Experience')
```

```
[27]: [Text(0.5, 1.0, 'Cluster vs Work Experience')]
```




0.0.6 Now take a look at the count of each cluster

```
[28]: automobile = automobile.fillna('Unknown')
      automobile["cluster"] = clusters
      automobile.head()
```


```
[28]:
```

	gender	married	age	graduated	profession	work_experience
0	Male	No	22	No	Healthcare	1.000000
1	Female	Yes	38	Yes	Engineer	2.619777
2	Female	Yes	67	Yes	Engineer	1.000000
3	Male	Yes	67	Yes	Lawyer	0.000000
4	Female	Yes	40	Yes	Entertainment	2.619777

	spending_score	family_size	var_1	cluster
0	Low	4.0	Cat_4	0
1	Average	3.0	Cat_4	3
2	Low	1.0	Cat_6	2
3	High	2.0	Cat_6	2
4	High	6.0	Cat_6	3

```
[29]: sns.catplot(x="cluster", kind="count", data=automobile)
```

```
[29]: <seaborn.axisgrid.FacetGrid at 0x7ffb588fad30>
```



The above barplots show the following insights: 1. Cluster 1: The average or typical customer in this cluster is middle-aged (around 40) with an average family size of 2.5 and a moderate ~ moderately high spending score. This customer group has the longest years of work_experience on average. This is the largest customer group. My recommendation to this company is to focus on utility — advertise versatile models that cater to a variety of driving conditions and preferences to maximize the reach to the largest pool of potential buyers.

2. Cluster 2: The typical customer in this cluster is a senior (over 70) with an average family size of 2 and a moderate ~ moderately high spending score. This is the smallest group and its members should be mostly retired.
3. Cluster 3: The typical customer is a slightly older middle-aged person with an average family size of 2.8 and the lowest spending score out of all groups with a less than 1.0 average. This group won't be able to afford expensive vehicles, probably because it has more family members to spend money on or this group as a whole do not spend beyond their needs. The company's strategy should be to advertise less expensive, mid-to-low tier vehicles.
4. Cluster 4: The typical customer is a young adult aged around 25 who has the highest spending score of all groups and the largest average family size! This is surprising because we normally assume older customers to have more spending power and larger family size, but for this

automobile company, this is simply not the case. This cluster is likely a group of young working professionals with high pay and/or heritage. The company should promote the most top-of-the-line car models to this group as their primary target consumer base because this group is most likely able to afford more expensive cars.

```
[30]: automobiles.columns
```

```
[30]: Index(['age', 'work_experience', 'spending_score', 'family_size',
          'profession_Artist', 'profession_Doctor', 'profession_Engineer',
          'profession_Entertainment', 'profession_Executive',
          'profession_Healthcare', 'profession_Homemaker', 'profession_Lawyer',
          'profession_Marketing', 'gender_Female', 'married_No', 'married_Yes',
          'graduated_No', 'graduated_Yes', 'var_1_Cat_1', 'var_1_Cat_2',
          'var_1_Cat_3', 'var_1_Cat_4', 'var_1_Cat_5', 'var_1_Cat_6',
          'var_1_Cat_7', 'label'],
          dtype='object')
```


```
[31]: automobiles.groupby("label")["graduated_Yes"].mean()
```

```
[31]: label
0      0.357340
1      0.771875
2      0.616933
3      0.740017
Name: graduated_Yes, dtype: float64
```

```
[32]: plt.figure(figsize=(20,5))
sns.catplot(x="profession", kind="count", hue='cluster', data=automobile,
            palette="ch:.25",
            height=5, aspect=3)
```

```
[32]: <seaborn.axisgrid.FacetGrid at 0x7ffb68ae3550>
```


<Figure size 1440x360 with 0 Axes>



We can see the following trends: 1. Artists, entertainment workers, and engineers are most represented in the 1st cluster. 2. Lawyers, artists, and executives are most represented in the 2nd cluster. 3. Artists, entertainment workers, and executives are most represented in the 3rd cluster. 4. Healthcare professionals, doctors, and artists are most represented in the 4th cluster.


```
[33]: sns.catplot(x="married", y="age", kind="bar", hue='cluster', data=automobile,
    ↪ palette="ch:.25",
    height=4, aspect=2)
```

```
[33]: <seaborn.axisgrid.FacetGrid at 0x7ffb8b9f3490>
```



```
[34]: sns.catplot(x="married", kind="count", hue='cluster', data=automobile,
    ↪ palette="ch:.25",
    height=4, aspect=2)
```

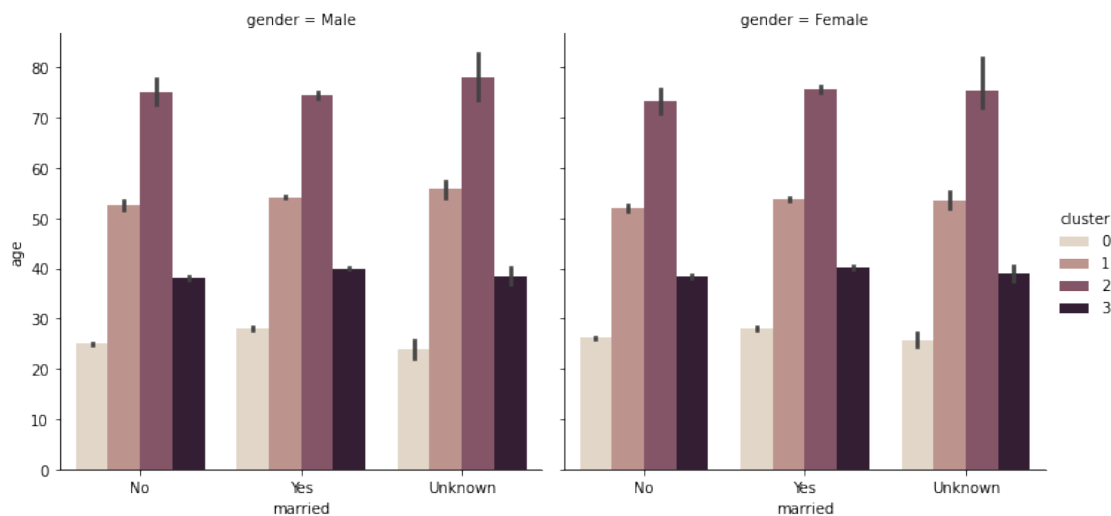
```
[34]: <seaborn.axisgrid.FacetGrid at 0x7ffb79c89370>
```



For all clusters except cluster 3 (actually the cluster 4 because python starts counting from 0), there are more married people than unmarried. The last cluster (represented by black) is the young professional group. It's understandable that most of them are unmarried.

```
[35]: sns.catplot(x="married", y="age", col="gender", kind="bar", hue='cluster',
↳data=automobile, palette="ch:.25")
```

```
[35]: <seaborn.axisgrid.FacetGrid at 0x7ffb79c89340>
```



The charts show that the age for each cluster for unmarried and married people is roughly equivalent. Age for female customers in each cluster is lower than the male counterparts.

```
[36]: automobile["gender"].value_counts()
```

```
[36]: Male      5841  
      Female    4854  
      Name: gender, dtype: int64
```

```
[37]: print(f"Male to Female Ratio is: {5841/4854}")
```

```
Male to Female Ratio is: 1.2033374536464772
```

0.0.7 This is the end of the analysis!